

3D Object Detection and Tracking Based on Streaming Data

Xusen Guo, Long Cheng and Kai Huang et al., Sun Yet-Sen University

Abstract—Recent approaches for 3D object detection have made tremendous progresses due to the development of deep learning. However, previous researches are mostly based on individual frames, leading to limited exploitation of information between frames. In this paper, we attempt to leverage the temporal information in streaming data and explore 3D streaming based object detection as well as tracking. Toward this goal, we set up a dual-way network for 3D object detection based on keyframes, then propagate predictions to non-key frames through a motion based interpolation algorithm guided by temporal information. Our framework is not only shown to have significant improvements on object detection compared with frame-by-frame paradigm, but also proven to produce competitive results on KITTI Object Tracking Benchmark, with 76.68% in MOTA and 81.65% in MOTP respectively.

I. INTRODUCTION

3D Object detection has received increasing attention during the last few decades due to the growing requirement for 3D information in robotics and autonomous driving, etc. Compared with 2D image, 3D data can provide accurate location of targets and characterize their shapes. Current approaches for 3D object detection are mostly carried out in three fronts: image based [1], [2], point clouds based [3], [4], [5], and multi-view fusion based [6], [7]. Most of these approaches have achieved competitive results but are limited to single frame input.

For most detection tasks, data are obtained in a streaming fashion, thus it is more natural to perform object detection based on streaming data. Compared with individual frames, consecutive frames can provide temporal information for long-term prediction of an object, which can reduce noisy detections over time. In addition, truncated and occluded targets can possibly be compensated by the information in its adjacent frames. Therefore, it is important to explore 3D object detection for streaming data.

Performing 3D object detection on streaming data is however complex. First of all, acquiring accurate 3D information in real world is hard. On one hand, camera data can provide rich appearance features but lack of depth information. On the other hand, though point cloud can precisely provide the position of objects, it is very sparse and thus difficult to represent their appearances. Secondly, how to correlate the temporal information in consecutive frames is not obvious. For example, generating 3D scene flow for temporal feature representation will need to determine the correspondence of thousands of points between frames, which is not straightforward and also challenging. Last but not least, due to the sheer numbers of frames in streaming data, frame-by-frame detection will introduce tremendous computational costs, which is impracticable for real-life applications.

This paper proposes a **Dual-way Object Detection and Tracking (DODT)** framework, as illustrated in *Figure 1*, to tackle aforementioned problems. We construct our framework based on the following observations: (1) Point clouds can be fused with images to obtain richer features for visual perception, which has been verified in [6], [7]. (2) Apart from 3D scene flow, temporal information can also be represented through cross-correlation between adjacent frames, as demonstrated in [8]. (3) Features along consecutive frames change gradually, thus we can process keyframes data only and propagate predictions to non-key frames. For (1), We simply utilize fusion scheme in [7] to aggregate features from different views. For (2), a *Temporal module* is introduced to capture temporal information through cross-correlation in BEV space, and then predicts object co-occurrence (whether an object occurs on two adjacent keyframes at the same time) and offsets between adjacent frames. Notably, different from [8], [9], our correlation operation is performed on proposal level, which is much more flexible and efficient. For (3), we design our framework into a dual-way fashion. Thus we can operate two adjacent keyframes simultaneously, and propagate prediction to other frames conveniently. Moreover, a specific network named *Shared RPN* is proposed to generate 3D proposals which shared by two detection branches.

To generate predictions for all frames, a motion based interpolation algorithm is developed to propagate keyframe predictions to non-key frames guided by the outputs of *Temporal module*. Meanwhile, multi-object tracking can also be accomplished through *tracking by detection* [10] paradigm. In summary, our contributions are threefold:

- We propose a dual-way framework named DODT, which performs 3D object detection and tracking precisely on streaming data based on keyframes.
- We propose *Temporal module* to encode temporal information across frames in proposal level, which is much more flexible and efficient than the methods in [8], [9].
- We develop a motion based interpolation algorithm for prediction propagation, obtaining significant performance improvements on both detection and tracking.

II. RELATED WORK

3D object detection. Currently, most approaches in 3D object detection are done in three fronts: image based, point cloud based, and multi-view fusion based. Image based approaches such as Mono3D [1] and 3DOP [2] use camera data only. Since images lack depth information, hand-crafted geometric features are required in these approaches. For the point cloud based methods, there are usually two ways: voxelization based and projection based, according to how point

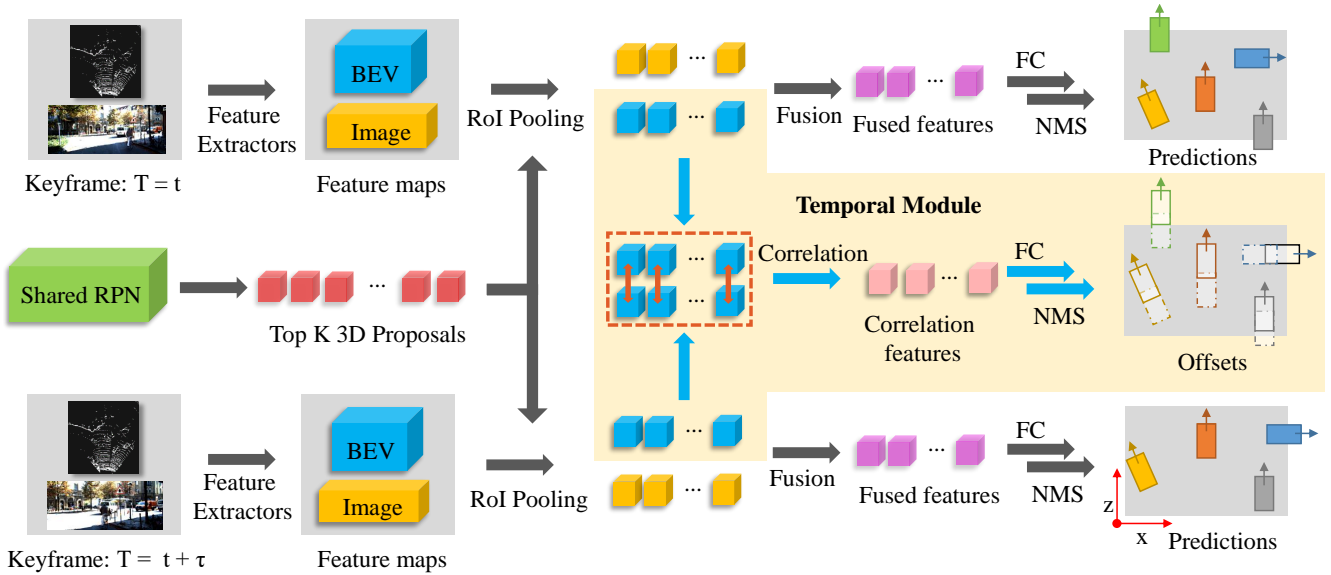


Fig. 1: DODT architecture. The structure in light yellow is *Temporal module*. Best viewed in color.

clouds features are represented. Voxelization based methods such as 3D FCN [11], Vote3Deep [12], VoxelNet [3], utilize a voxel grid to encode features. These approaches suffer from the sparsity of point clouds and enormous computation costs in 3D convolution. Projection based methods such as PIXOR [4], Complex-YOLO [5], [13], attempt to project point clouds to a perspective view (e.g. bird eye view) and apply image-based feature extraction techniques. However, due to the sparsity of point clouds, features after projection are usually insufficient for accurate object detection, especially for small targets. Multi-view fusion based approaches such as F-PointNet [14], MV3D [6], AVOD [7], try to fuse point cloud features with image features for accurate object detection. These methods distinguish from each other mainly on how data is fused. Our detection branches are constructed based on AVOD largely, but RPN is replaced with *Shared RPN*, and temporal features are integrated with dual-way structure.

Video object detection. Nearly all existing methods in video object detection incorporate temporal information on either feature level or final box level. FGFA [15] leverages temporal coherence on feature level, it warps the nearby frames feature maps to a reference frame for feature enhancement according to flow emotion. On the other hand, T-CNN [16], [17] leverages precomputed optical flows to propagate predicted bounding boxes to neighboring frames; Seq-NMS [18] improves NMS algorithm for video by constructing sequences along nearby high-confidence bounding boxes from consecutive frames. They all exploit temporal information in final box level. There are also a few approaches trying to learn temporal features between consecutive frames without optical flow. D&T [8] proposes a two branches detection network for object detection and tracking simultaneously in video. The network learns temporal information representation through computing convolutional cross-correlation between frames. Our DODT approach is mainly inspired by D&T. However, we develop this idea to 3D space and restrain correlation operation in proposal level instead of a fixed window. Moreover, data association with two adjacent

frames often suffers from drift and loss of targets in one frame, etc. D&T ignores these problems, while our approach can handle these issues well by performing motion based interpolation algorithm.

3D multi-object tracking. Existing 3D multi-object tracking approaches are mostly implemented based on tracking by detection. For example, FaF [19] jointly reasons about 3D detection, tracking and motion forecasting taking a 4D tensor created from multiple consecutive frames. It can aggregate the detection information for the past n timestamps to produce accurate tracklets. 3D-CNN/PMBM [20] trains a DNN to detect and estimate the distance to objects from a single image, and then feeds the detections to a Poisson multi-Bernoulli mixture tracking filter for 3D tracking. DSM [21] first predicts 3D bounding boxes in continuous frames and then associates detections using a *Matching net* and a *Scoring net*, which is similar to our approach. However, their 3D detector is single frame based approach MV3D [6], and temporal features between frames are mostly ignored. Moreover, their bounding boxes association is done by solving a linear program using a network, while ours is done simply by applying an IoU based algorithm [22].

III. METHODOLOGY

In this section, we first give an overview of the network (Sec. A) that performs 3D object detection given two adjacent keyframes as inputs. Then we introduce the *shared RPN* (Sec. B) that generates 3D proposals shared by two detection branches. Sec. C shows how *Temporal module* encodes cross-correlation features and then predicts the co-occurrence and offsets of a target in two adjacent keyframes. Sec. D shows how we implement motion based interpolation algorithm and accomplish streaming level detection and multi-object tracking simultaneously.

A. DODT Model Structure

We aim at performing 3D object detection and tracking based on streaming data. To this end, we design DODT into a dual-way network. *Figure 1* illustrates the whole pipeline.

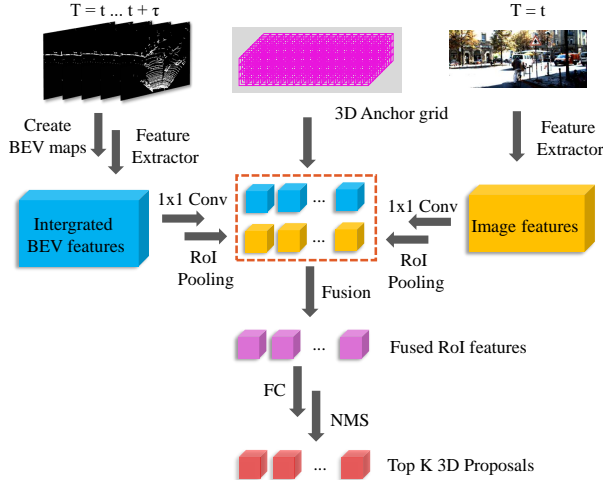


Fig. 2: Shared RPN module.

With the help of the dual-way structure, the network can be fed with two adjacent keyframes data simultaneously. Consisting of an image and point cloud BEV maps (following the procedure described in MV3D [6]), the keyframe data are first fed to feature extractors to get corresponding feature maps. Feature extractors are designed following the procedure described in AVOD [7], which are constructed in an encoder-decoder fashion resulting in a full resolution feature map. To extract feature crops from image and BEV feature maps, we again follow the idea proposed in AVOD. Given a 3D anchor generated by RPN, two view specific ROIs are obtained by projecting the anchor onto the image and BEV feature maps, then the corresponding feature crops are extracted and RoI pooling is performed individually. After that, a *early fusion* scheme from MV3D [6] is used to combine multi-view features in proposal level. Finally, by applying two *fc* networks for classification and box regression and following 2D non-maximum suppression (NMS), the final predictions are produced. Meanwhile, fed with BEV feature crops of two branches, *Temporal module* performs correlation operation on feature crop pairs to produce cross-correlation features. The cross-correlation features are then used to predict object co-occurrence and offsets between two keyframes for further streaming-level detection and tracking.

The whole network is designed in an end-to-end form. Our multi-task objective is represented as

$$L = w_{cls}L_{cls} + w_{reg}L_{reg} + w_{co}L_{co} + w_{corr}L_{corr} \quad (1)$$

where L_{cls} , L_{co} are cross-entropy loss for classification and object co-occurrence, and L_{reg} , L_{corr} are smooth *LI* loss for box regression and offsets regression. w_{cls} , w_{reg} , w_{co} , w_{corr} are loss weights with values 1.0, 5.0, 1.0, 1.0, respectively. During training, L_{reg} and L_{corr} are normalized by the number of proposals, while L_{cls} , L_{co} are normalized by the number of positive proposals.

B. Shared RPN

We transform RPN network in AVOD [7] to our *shared RPN* (as shown in Figure 2), which can produce 3D proposals shared by both detection branches. To ensure proposals generated by *shared RPN* are suitable for both keyframes, we

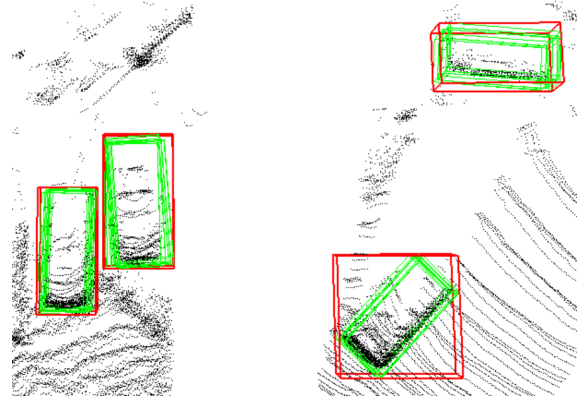


Fig. 3: Five consecutive point clouds in the same coordinate system. Green boxes illustrate default labels of five frames, while red boxes are the new axis-aligned labels for RPN training.

create integrated BEV feature maps based on frames between two adjacent keyframes, inclusive. Note that point cloud shows accurate 3D localization of objects, we can simply transform these frames to the same coordinate system and then integrate them. Since point cloud is extremely sparse and features are encoded by projection, this process does not increase any computational cost but enhances the BEV feature maps greatly. Assume there are five frames in a frame segment, due to object movement, the locations of the same objects in five frames are shift. For training convenience, we replace original proposal boxes labels with new axis-aligned labels that could cover all bounding boxes of an object in five frames. Figure 3 illustrates these two kinds of boxes. Though this process enlarges proposals slightly, it can ensure that new labels are suitable for proposal prediction in both detection branches. In addition, to obtain accurate proposals, we fuse BEV feature maps with image data to enhance object appearance features. Since one image contains enough features in a short time, our module only aggregates the features in first image of a frame segment. Subsequent processes are similar with RPN module described in AVOD, we refer the reader to [7] for more information.

C. Temporal Module

Our *Temporal Module* is demonstrated in the light yellow area in Figure 1. Given two sets of point cloud BEV features crops $F_t, F_{t+\tau}$, a set of cross frame feature pairs can be constructed as $\{(F_t^i, F_{t+\tau}^i) \mid i \in \{0, 1, \dots, N\}\}$, where $F_t^i, F_{t+\tau}^i$ are features extracted by i -th proposal from frame t and frame $t+\tau$ respectively, τ is temporal stride and N is the number of 3D proposals. Note that proposals generated by our RPN are shared by two detection branches, thus the set can be obtained conveniently. After the correspondence of feature crops is constructed, correlation operation can be performed on feature pair $(F_t^i, F_{t+\tau}^i)$ to compute cross-correlation features over frames. The features are then fed to a classification head for object co-occurrence prediction and a regression head for box offsets prediction respectively. Object co-occurrence shows the probability of an object exists on both keyframes, it provides an extra guarantee for successful implementation of the interpolation algorithm. When an object is lost in one of the two keyframes, object co-

occurrence determines whether a track birth or death occurs, or just a mis-detection in one keyframes. More details are available in Sec. D.

For a target we have ground truth $d^t = (p_{co}, d_x^t, d_z^t, d_{ry}^t)$ for frame t in *Temporal Module*, and similarly $d^{t+\tau}$ for frame $t+\tau$, denoting object co-occurrence probability, its horizontal and vertical center coordinates as well as orientation in BEV map. p_{co} is 1.0 if the target exists on both keyframes, 0 otherwise. The box offsets label $\delta^{t,t+\tau} = (\delta_x^{t,t+\tau}, \delta_z^{t,t+\tau}, \delta_{ry}^{t,t+\tau})$ is then represented as

$$\delta^{t,t+\tau} = \begin{cases} (\frac{d_x^{t+\tau}-d_x^t}{d_w^t}, \frac{d_z^{t+\tau}-d_z^t}{d_l^t}, \frac{d_{ry}^{t+\tau}-d_{ry}^t}{d_{ry}^t}) & p_{co} = 1.0 \\ (0.0, 0.0, 0.0) & otherwise \end{cases} \quad (2)$$

where d_w^t, d_l^t are width and height of the object.

D. 3D Streaming Object Detection and Tracking

Algo. 1 shows how we perform streaming level detection and tracking with DODT outputs. Fed with two detections lists $D^t, D^{t+\tau}$ and a offsets list $\Delta^{t,t+\tau}$ of adjacent keyframes, the algorithm first utilizes linear interpolation to calculate detections in intermediate frames if an object exists on both keyframes. Detections association between two keyframes is completed using *getMatched* function. Given a rectified detection $d_i^t + \Delta d_i^{t,t+\tau}$ ($\Delta d_i^{t,t+\tau}$ is decoded from $\delta_i^{t,t+\tau}$) in first keyframe, *getMatched* function performs a 3D IoU based algorithm with a threshold 0.2 to select the best matched detection in second keyframe. As for the detection that exists only in one keyframe, the match fails and the function outputs *None*. In this case, if its co-occurrence probability p_{co}^i is not less than p_{co}^{max} (0.5 in our experiment), we regard it as a mis-detection in another keyframe, then we utilize offsets predictions to compute corresponding detection in another keyframe and do interpolation subsequently. Otherwise we regard it as a track birth or death, then a motion model is leveraged to propagate the detections.

The Motion model holds the hypothesis that the target velocity is constant within a short time period, independent of camera ego-motion. Thus we can calculate target state in next frame using historical frames states and the increments. In detail, we maintain a global velocity $v = (v_x, v_z, v_{ry})$ in BEV space, which will be updated to $\alpha v + (1 - \alpha) \frac{\delta^{t,t+\tau}}{\tau} (d_w^t, d_l^t, d_{ry}^t)$ ($\alpha = 0.8$ in our experiment) once a new detection is matched with the track. As for the end of a track, since there is no way for us to locate its exactly end frame with only keyframes predictions, we extend the track by 3 frames according to the trend of movement instead (with the maintained global velocity). If the track is beyond the range of BEV feature maps, the extension will be terminated earlier. Similarly when it turns to the start of a track.

In addition, we observe that a track sometimes suffers from opposite orientation in nearby keyframes, which will lead to inferior performance. To address this issue, we operate a simple orientation correction algorithm during box interpolation and track extension. Consider a track T_{match}^i and its next matched detection D_{match}^i , if the difference of orientation in T_{match}^i and D_{match}^i is grater than $\frac{\pi}{2}$, we add

a π to the orientation in D_{match}^i so that its orientation can be roughly consistent with the track.

Multi-object tracking can be accomplished with streaming level object detection simultaneously. Note that after performing interpolation algorithm on keyframe pairs, data association between them is also obtained. We only need to associate detections between different keyframes pairs to link tracklets over time and build long-term object tubes. Since our method associates a frame segment to the track at one time, this shows an near online tracking approach.

Algorithm 1: Motion based Interpolation Algorithm

```

1 Input:
    $D^t = [d_0^t, d_1^t, \dots, d_{N_t}^t], D^{t+\tau} = [d_0^{t+\tau}, d_1^{t+\tau}, \dots, d_{N_{t+\tau}}^{t+\tau}],$ 
    $\Delta^{t,t+\tau} = [\delta_0^{t,t+\tau}, \delta_1^{t,t+\tau}, \dots, \delta_N^{t,t+\tau}], N = \max\{N_t, N_{t+\tau}\}$ 
2 Output:  $D = [D^t, D^{t+1}, \dots, D^{t+\tau}]$ 
3 Initialize:  $D_{temp} = D^{t+\tau}, D, p_{co}^{max}$ 
4 for  $d_i^t \in D^t$  do
5    $\Delta d_i^{t,t+\tau} =$ 
      $(d_{i,w}^t \cdot \delta_{i,x}^{t,t+\tau}, 0, d_{i,l}^t \cdot \delta_{i,z}^{t,t+\tau}, 0, 0, d_{i,ry}^t \cdot \delta_{i,ry}^{t,t+\tau})$ 
      $d' = \text{getMatched}(d_i^t + \Delta d_i^{t,t+\tau}, D_{temp})$ 
6   if  $d' \neq \text{None}$  then
7      $d_i^{t+1}, \dots, d_i^{t+\tau-1} = \text{Interpolate}(d_i^t, d')$ 
8     remove  $d'$  from  $D_{temp}$ 
9   else if  $p_{co}^i \geq p_{co}^{max}$  then
10     $d_i^{t+1}, \dots, d_i^{t+\tau} = \text{Interpolate}(d_i^t, d_i^t + \Delta d_i^{t,t+\tau})$ 
11   else
12    generate  $(d_i^{t+1}, \dots, d_i^{t+\tau-1})$  by motion model
13 if  $D_{temp}$  is not empty then
14   for  $d_j^{t+\tau} \in D_{temp}$  do
15     if  $p_{co}^j > p_{co}^{max}$  then
16        $d_j^t, \dots, d_j^{t+\tau-1} =$ 
          $\text{Interpolate}(d_j^{t+\tau} - \Delta d_j^{t,t+\tau}, d_j^{t+\tau})$ 
17     else
18       generate  $(d_j^{t+1}, \dots, d_j^{t+\tau-1})$  by motion model

```

IV. EXPERIMENTS

A. Datasets and Training

Datasets and Preprocessing. We use KITTI object tracking Benchmark [23] for evaluation. It consists of 21 training sequences and 29 test sequences with vehicles annotated in 3D. For object detection, we split 21 training sequences into two parts according to their sequence number, odd numbered sequences for training and the rest for validation. For tracking evaluation, we train our model in all 21 training sequences and evaluate on test datasets. Similar to the data preprocessing in AVOD [7], we crop point clouds at $[-40, 40] \times [0, 70] \times [0, 2.5]$ meters along X, Z, Y axis respectively to contain points within the field of camera view. To make the system invariant to the speed of the ego-car, we calculate the displacement of the observer between adjacent keyframes using IMU data, and then transform the coordinates in second keyframe accordingly.

Methods	Modules	IoU = 0.5			IoU = 0.7			FPS
		Easy	Moderate	Hard	Easy	Moderate	Hard	
AVOD[7]	-	90.13 / 90.91	80.00 / 81.79	71.61 / 81.79	76.00 / 90.90	57.23 / 81.73	56.13 / 72.69	10.0
DODT($\tau = 1$)	-	88.28 / 99.97	85.74 / 90.90	86.14 / 90.89	83.44 / 90.82	67.48 / 90.79	61.24 / 90.80	6.7
DODT($\tau = 1$)	T	88.32 / 99.99	86.53 / 90.90	86.71 / 90.90	83.60 / 90.82	68.93 / 90.80	62.69 / 90.81	5.9
DODT($\tau = 1$)	M	89.99 / 99.95	87.86 / 90.87	87.81 / 90.86	86.89 / 90.89	73.96 / 90.83	67.07 / 81.79	6.5
DODT($\tau = 1$)	T + M	90.63 / 99.95	89.07 / 90.90	88.79 / 90.90	88.74 / 90.91	75.27 / 90.84	68.75 / 90.57	5.7
DODT($\tau = 2$)	T + M	90.60 / 99.94	89.19 / 90.91	88.91 / 90.88	88.90 / 90.92	76.64 / 90.85	75.81 / 90.83	8.6
DODT($\tau = 3$)	T + M	90.61 / 99.98	89.01 / 90.89	88.84 / 90.89	88.81 / 90.91	76.38 / 90.86	75.83 / 90.85	11.4
DODT($\tau = 4$)	T + M	90.55 / 99.94	88.82 / 90.88	88.34 / 90.87	88.43 / 90.91	75.70 / 90.82	68.75 / 90.82	14.3
DODT($\tau = 5$)	T + M	87.98 / 90.91	85.57 / 90.87	86.01 / 90.87	81.59 / 90.81	67.30 / 90.76	61.35 / 81.73	17.1
DODT($\tau = 6$)	T + M	78.77 / 90.75	70.88 / 90.71	71.65 / 81.70	71.71 / 90.44	55.86 / 81.50	56.80 / 81.51	20.0

TABLE I: We report AP_{3D}/AP_{BEV} (in %) of the **Car** category on validation datasets, corresponding to average precision of 3D object detection and the bird’s eye view. T is *Temporal Module*, M is *Motion based interpolation algorithm*. τ is temporal stride.

We notice that the labels provided in KITTI object tracking datasets are incomplete. For example, first row in Figure 4 illustrates several frames and labels in sequence 0. We can see that some objects (marked by red rectangles) are not labeled, even though they can be well observed in frame 118 and 120 and are labeled in frame 128. However, our model can well predict these objects, as shown in second row. In order to evaluate our model as accurately as possible, we add these missing labels manually.

Training and testing. We train our network only for *Car* category temporarily, following most of the super-parameter settings in AVOD [7] during training and testing. The network is trained for 120K iterations with a batch size 1, using an ADAM [24] optimizer with an initial learning rate of 0.0001 that is decayed exponentially every 30K iterations with a decay factor of 0.8. During proposal generation, anchors with IoU less than 0.3 are considered background and greater than 0.5 are objects. To remove redundant proposals, 2D NMS is performed at an IoU threshold of 0.8 in BEV to keep the top 1024 proposals during training, while at inference time, the top 300 proposals are kept.

B. Results

Shared RPN. To evaluate the performance of our *Shared RPN*, we implement a non-shared version of RPN. It predicts proposals for each keyframe independently based on each feature maps. A comparison of proposal prediction accuracy between *Non-shared RPN* and *Shared RPN* is shown in Figure II. Results show that *Shared RPN* outperforms *Non-shared RPN* by 0.66%, which indicates that the shared mechanism in RPN promotes the accuracy of proposal prediction.

Method	Non-shared RPN	Shared RPN
Accuracy(%)	97.81	98.47

TABLE II: Comparison of proposal prediction accuracy.

Streaming level detection. The main results on 3D object detection are summarized in Table I. We first evaluate the effectiveness of the dual-way structure, *Temporal Module* and the interpolation algorithm in 3D object detection with a temporal stride $\tau = 1$. Several important trends can be observed: **1)** Compared with original AVOD [7] model trained on the same datasets, our base model (without *Temporal Module* and interpolation algorithm) shows improvements in

all settings with IoU = 0.7. These improvements indicate that the dual-way structure with *Shared RPN* contribute to detection performance, as it can aggregate features in consecutive frames to produce more accurate proposals. **2)** The introduction of *Temporal Module* and interpolation algorithm is conducive to model performance. Compared with our base model, adding *Temporal Module* brings 0.16%, 1.45% and 1.45% gain in *Easy*, *Moderate* and *Hard* setting respectively with a overlap of 0.7. It shows that the aid of object co-occurrences information contributes to the detection in occluded and truncated objects more than normal ones. The comparison also shows that our interpolation algorithm improves model accuracy by 3.45%, 6.48% and 5.83% in three setting respectively with a overlap of 0.7. These improvements indicate that our algorithm works well in all settings. It’s not surprising because the algorithm has the ability to reject false positive predictions and generate new results by extending the trajectories. **3)** *Temporal Module* and interpolation algorithm can work synergistically and improve model accuracy by 1-2% additionally in all setting, since the interpolation algorithm can work better with the results of *Temporal Module* (see Sec. D in methodology for details).

Secondly, we investigate the effect of temporal stride τ on model performance. We train six models with $\tau = \{1, 2, 3, 4, 5, 6\}$ respectively and then perform the interpolation algorithm to generate predictions for all frames. Results are shown in Table I. DODT ($\tau = 2$) shows the best result in *Easy* and *Moderate* settings with IoU = 0.7, by 88.90% and 76.64% in AP_{3D} respectively. DODT ($\tau = 3$) performs the best in *Hard* setting, by 75.84% in AP_{3D} . Compared with DODT ($\tau = 1$), model performance with $\tau = 2, 3$ are boosted in hard objects, by about 1% in *Moderate* setting and about 7% in *Hard* setting with IoU = 0.7. These gains demonstrate that the detection of truncated and occluded targets can benefit greatly from a large temporal stride. Bottom two rows in Figure 4 shows how our interpolation algorithm improves model performance. Boxes in magenta are DODT predictions with $\tau = 3$ while yellow with $\tau = 1$. The results shows that a larger τ can reduce noisy detections and generate more true positive predictions. However, Table I also shows that a too large τ leads to a significant decay of accuracy. As temporal stride increasing, there are more tracks start or end within two adjacent keyframes, thus our interpolation algorithm would fail to generate most objects

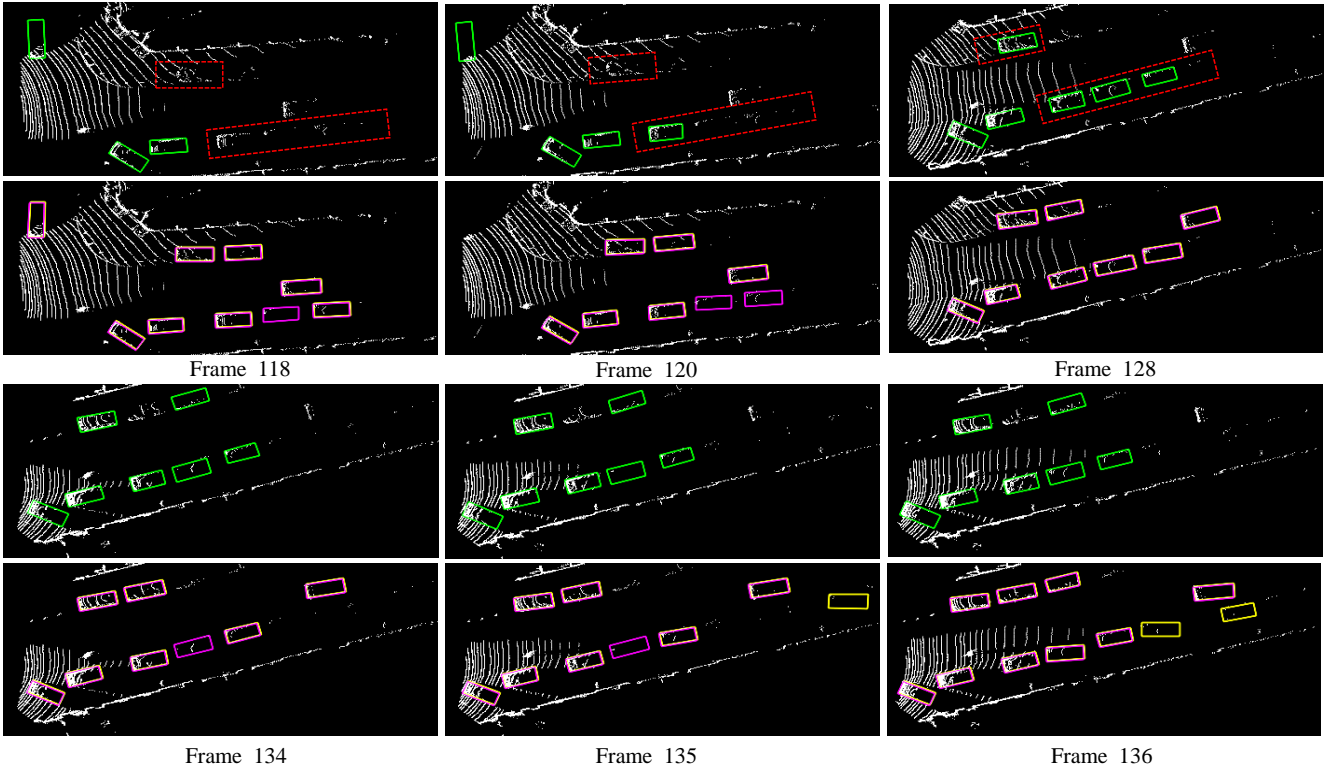


Fig. 4: Visualization of labels and our predictions in sequences 0. Boxes in **green** are official ground truth, boxes in **yellow** are results predicted with temporal stride $\tau = 1$, and boxes in **magenta** are results predicted with temporal stride $\tau = 3$. Boxes in mixed color are magenta boxes overlapped with yellow boxes. Best viewed in color.

Methods	Modules	MOTA(%) \uparrow	MOTP(%) \uparrow	MT(%) \uparrow	ML(%) \downarrow	IDS \downarrow	FRAG \downarrow
AVOD[7]	-	66.05	82.97	46.22	12.18	2	113
DDOT($\tau = 3$)	-	76.53	83.93	68.91	7.14	32	80
DDOT($\tau = 3$)	T	77.52	83.75	69.33	7.56	37	77
DDOT($\tau = 3$)	M	78.73	83.93	68.49	9.55	2	48
DDOT($\tau = 3$)	T+M	79.72	83.55	71.85	5.46	7	66

TABLE III: Ablation study on KITTI Tracking validation datasets.

Method	MOTA(%) \uparrow	MOTP(%) \uparrow	MT(%) \uparrow	ML(%) \downarrow	IDS \downarrow	FRAG \downarrow	FPS \uparrow
Complexer-YOLO[13]	75.70	78.46	58.00	5.08	1186	2096	100.0
DSM[21]	76.15	83.42	60.00	8.31	296	868	10.0 (GPU)
3D-CNN/PMBM[20]	80.39	81.26	62.77	6.15	121	613	71.4
3DT[25]	84.52	85.64	73.38	2.77	377	847	33.3
DDOT(ours)	76.68	81.65	60.77	11.69	63	384	76.9

TABLE IV: Comparison of publicly available methods of 3D multi-object tracking in the KITTI Tracking Benchmark. The time for object detection is not included in the specified runtime.

near the start or end of a track. Moreover, temporal information in a longer time period is more difficult to represent.

We also compute the inference time in streaming level detection. The total runtime of two detection branches is 175ms on a Tesla P100 GPU. As temporal stride τ increasing, the time cost for detection is unchanged, and interpolation time cost increases negligibly, thus the total inference time is almost unchanged, but FPS increases by multiple. We choose $\tau = 3$ for our following experiment, which is a good trade-off between speed and accuracy.

Multi-object tracking. For multi-object tracking, We investigate the contribution of different modules to tracking accuracy in KITTI tracking validation datasets. Results in Table III shows that 1) our approach outperforms original AVOD model by a large margin (e.g. MOTA by 13.67%, MOTP by 0.58%, MT by 25.63%, etc.). These improvements mainly benefit from better object detection results. 2) Compared with base model, performances of the model aided with

two modules are superior in nearly all tracking metrics, as *Temporal Module* makes the data association more accurate, and box interpolation helps to complete both ends of tracks.

We also compare our approach to publicly available methods of multiple object tracking in 3D on KITTI Tracking Benchmark. Results are shown in Table IV. We can see that our approach is competitive with the state-of-the-art in some metrics (e.g. IDS and FRAG). For multiple object tracking accuracy, our method outperforms Complexer-YOLO [13] and DSM [21], but behind 3D-CNN/PMBM [20] and 3DT [25]; For trajectory ID-switches and fragmentations, our method outperforms all other methods. Note that 3D-CNN/PMBM utilizes PMBM filter for data association while ours is a simple IoU based method, and 3DT trained their system additionally on new datasets collected on a realistic 3D virtual environments. Moreover, labels of test datasets may be incomplete just like the case in training datasets, thus the tracking performance of our approach may be better than current ones. Also, our approach shows a competitive results in runtime with 76.9 FPS (on a 3.2 GHz Intel CPU). Note that we could not separate predictions propagation from data association in Algo. 1 easily, thus the tracking runtime also includes box interpolation process.

V. CONCLUSIONS

We propose DDOT, a unified framework for simultaneous 3D object detection and tracking based on streaming data. The network is a dual-way structure and can process two keyframes at the same time. Embedded with a *Temporal Module* to encode the temporal information among adjacent

keyframes and a motion based interpolation algorithm to propagate predictions to non-key frames, our network can perform object detection and tracking in a very efficient way.

REFERENCES

- [1] X. Chen, K. Kundu, Z. Zhang, H. Ma, S. Fidler, and R. Urtasun, "Monocular 3d object detection for autonomous driving," in *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2016, pp. 2147–2156.
- [2] X. Chen, K. Kundu, Y. Zhu, H. Ma, S. Fidler, and R. Urtasun, "3d object proposals using stereo imagery for accurate object class detection," *IEEE transactions on pattern analysis and machine intelligence*, vol. 40, no. 5, pp. 1259–1272, 2018.
- [3] Y. Zhou and O. Tuzel, "Voxelnet: End-to-end learning for point cloud based 3d object detection," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2018, pp. 4490–4499.
- [4] B. Yang, W. Luo, and R. Urtasun, "Pixor: Real-time 3d object detection from point clouds," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2018, pp. 7652–7660.
- [5] M. Simon, S. Milz, K. Amende, and H.-M. Gross, "Complex-yolo: An euler-region-proposal for real-time 3d object detection on point clouds," in *European Conference on Computer Vision*. Springer, 2018, pp. 197–209.
- [6] X. Chen, H. Ma, J. Wan, B. Li, and T. Xia, "Multi-view 3d object detection network for autonomous driving," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2017, pp. 1907–1915.
- [7] J. Ku, M. Mozifian, J. Lee, A. Harakeh, and S. L. Waslander, "Joint 3d proposal generation and object detection from view aggregation," in *2018 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. IEEE, 2018, pp. 1–8.
- [8] C. Feichtenhofer, A. Pinz, and A. Zisserman, "Detect to track and track to detect," in *Proceedings of the IEEE International Conference on Computer Vision*, 2017, pp. 3038–3046.
- [9] A. Dosovitskiy, P. Fischer, E. Ilg, P. Hausser, C. Hazirbas, V. Golkov, P. Van Der Smagt, D. Cremers, and T. Brox, "FlowNet: Learning optical flow with convolutional networks," in *Proceedings of the IEEE international conference on computer vision*, 2015, pp. 2758–2766.
- [10] P. Lenz, A. Geiger, and R. Urtasun, "Followme: Efficient online min-cost flow tracking with bounded memory and computation," in *Proceedings of the IEEE International Conference on Computer Vision*, 2015, pp. 4364–4372.
- [11] B. Li, "3d fully convolutional network for vehicle detection in point cloud," in *2017 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. IEEE, 2017, pp. 1513–1518.
- [12] M. Engelcke, D. Rao, D. Z. Wang, C. H. Tong, and I. Posner, "Vote3deep: Fast object detection in 3d point clouds using efficient convolutional neural networks," in *2017 IEEE International Conference on Robotics and Automation (ICRA)*. IEEE, 2017, pp. 1355–1361.
- [13] M. Simon, K. Amende, A. Kraus, J. Honer, T. Samann, H. Kaulbersch, S. Milz, and H. Michael Gross, "Complexer-yolo: Real-time 3d object detection and tracking on semantic point clouds," in *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR) Workshops*, June 2019.
- [14] C. R. Qi, W. Liu, C. Wu, H. Su, and L. J. Guibas, "Frustum pointnets for 3d object detection from rgb-d data," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2018, pp. 918–927.
- [15] X. Zhu, Y. Wang, J. Dai, L. Yuan, and Y. Wei, "Flow-guided feature aggregation for video object detection," in *Proceedings of the IEEE International Conference on Computer Vision*, 2017, pp. 408–417.
- [16] K. Kang, H. Li, J. Yan, X. Zeng, B. Yang, T. Xiao, C. Zhang, Z. Wang, R. Wang, X. Wang *et al.*, "T-cnn: Tubelets with convolutional neural networks for object detection from videos," *IEEE Transactions on Circuits and Systems for Video Technology*, vol. 28, no. 10, pp. 2896–2907, 2018.
- [17] K. Kang, W. Ouyang, H. Li, and X. Wang, "Object detection from video tubelets with convolutional neural networks," in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2016, pp. 817–825.
- [18] W. Luo, B. Yang, and R. Urtasun, "Fast and furious: Real time end-to-end 3d detection, tracking and motion forecasting with a single convolutional net," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2018, pp. 3569–3577.
- [19] S. Scheidegger, J. Benjaminsson, E. Rosenberg, A. Krishnan, and K. Granström, "Mono-camera 3d multi-object tracking using deep learning detections and pmhm filtering," in *2018 IEEE Intelligent Vehicles Symposium (IV)*. IEEE, 2018, pp. 433–440.
- [20] D. Frossard and R. Urtasun, "End-to-end learning of multi-sensor 3d tracking by detection," in *2018 IEEE International Conference on Robotics and Automation (ICRA)*. IEEE, 2018, pp. 635–642.
- [21] E. Bochinski, T. Senst, and T. Sikora, "Extending iou based multi-object tracking by visual information," *AVSS. IEEE*, 2018.
- [22] A. Geiger, P. Lenz, C. Stiller, and R. Urtasun, "Vision meets robotics: The kitti dataset," *The International Journal of Robotics Research*, vol. 32, no. 11, pp. 1231–1237, 2013.
- [23] D. P. Kingma and J. Ba, "Adam: A method for stochastic optimization," *arXiv preprint arXiv:1412.6980*, 2014.
- [24] H.-N. Hu, Q.-Z. Cai, D. Wang, J. Lin, M. Sun, P. Krhenbhl, T. Darrell, and F. Yu, "Joint monocular 3d detection and tracking," 2019.